

Problem 1

(a) (a)

Algorithm 1: Loading, Preprocessing, Reshaping Data

```
import numpy as np
import keras
from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(f'Orig_train_set_range: [{np.min(x_train)}, {np.max(x_train)}]')
print(f'Orig_test_set_range: [{np.min(x_test)}, {np.max(x_test)}]')
"""
Orig train set range: [0,255]
Orig test set range: [0,255]
"""

x_train = np.divide(x_train, 255)
x_test = np.divide(x_test, 255)

print(f'Normalized_train_set_range: [{np.min(x_train)}, {np.max(x_train)}]')
print(f'Normalized_test_set_range: [{np.min(x_test)}, {np.max(x_test)}]')
"""
Normalized train set range: [0.0,1.0]
Normalized test set range: [0.0,1.0]
"""

print('Orig_train_samples_shape:', x_train.shape)
print('Orig_test_samples_shape:', x_test.shape)
"""
Orig train samples shape: (60000, 28, 28)
Orig test samples shape: (10000, 28, 28)
"""

x_train = np.reshape(x_train, (x_train.shape[0], 784))
x_test = np.reshape(x_test, (x_test.shape[0], 784))

print('Vectorized_train_samples_shape:', x_train.shape)
print('Vectorized_test_samples_shape:', x_test.shape)
"""
Vectorized train samples shape: (60000, 784)
Vectorized test samples shape: (10000, 784)
"""

one_hot_train = keras.utils.to_categorical(y_train, num_classes=10)
one_hot_test = keras.utils.to_categorical(y_test, num_classes=10)
```

(b)

Algorithm 2: Plot Some Data

```
import numpy as np
```

```
import matplotlib.pyplot as plt

def plot_5(trn, labs):
    f, axrow = plt.subplots(1, 5)
    axrow[0].axis('off')
    axrow[0].imshow(trn[0], cmap='gray')
    axrow[0].set_title(f"class={labs[0]}")
    axrow[1].axis('off')
    axrow[1].imshow(trn[1], cmap='gray')
    axrow[1].set_title(f"class={labs[1]}")
    axrow[2].axis('off')
    axrow[2].imshow(trn[2], cmap='gray')
    axrow[2].set_title(f"class={labs[2]}")
    axrow[3].axis('off')
    axrow[3].imshow(trn[3], cmap='gray')
    axrow[3].set_title(f"class={labs[3]}")
    axrow[4].axis('off')
    axrow[4].imshow(trn[4], cmap='gray')
    axrow[4].set_title(f"class={labs[4]}")
    # plt.show()
    plt.savefig('./first_5.png')
    return
```

```
plot_5(x_train, y_train) # See Figure 1
```



Figure 1: Plot of the first 5 training examples

(c)

Algorithm 3: Build A Model

```
def build_model(dp=False, sigma=0.01):
    units=512
    hiddens=2
    model = Sequential()
    model.add(Dense(units, input_dim=(28*28), activation='relu'))
    model.add(Dense(units, activation='relu'))
    model.add(Dense(10, activation='softmax'))
    if dp: optim = DPSGD(sigma=sigma)
    else: optim = 'rmsprop'
    model.compile(optimizer=optim, loss='categorical_crossentropy', metrics=
```

```
plot_training([acc],[loss], 'non_priv')
```




Figure 2: Non-Private Training Accuracy Loss, 20 Epochs

Problem 2

Algorithm 6: Plot Training Accuracy Loss

```
class DPSPGD(Optimizer):
    """
    Differentially Private Stochastic gradient descent optimizer.
    # Arguments
    lr: float >= 0. Learning rate.
    decay: float >= 0. Learning rate decay over each update.
    sigma: float >= 0. scale parameter of (Gaussian) noise distribution
    """
    def __init__(self, lr=0.01, decay=0., sigma=0.0001, **kwargs):
        super(DPSPGD, self).__init__(**kwargs)
        with K.name_scope(self.__class__.__name__):
            self.iterations = K.variable(0, dtype='int64',
                                         name='iterations')
            self.lr = K.variable(lr, name='lr')
            self.initial_decay = decay
            self.sigma = sigma
        @interfaces.legacy_get_updates_support
        def get_updates(self, loss, params):
            """ Noise injected in lines below """
            grads=self.get_gradients(loss, params)
            noisy_grads = [np.add(g,
                                  keras.backend.random_normal(g.shape, mean=0.,
                                                                stddev=self.sigma**2)) for g in grads]
            self.updates = [K.update_add(self.iterations, 1)]
            lr = self.lr
            if self.initial_decay > 0:
                lr = lr * (1. / (1. + self.decay *
                                K.cast(self.iterations,
```

```

        K.dtype(self.decay))))
    shapes = [K.int_shape(p) for p in params]
        # returns array of tensors the size of each param
    self.weights = [self.iterations]
    for p, g in zip(params, noisy_grads):
        v = -lr * g # velocity
        """ update occurs here """
        new_p = p + v
        if getattr(p, 'constraint', None) is not None:
            new_p = p.constraint(new_p)
        self.updates.append(K.update(p, new_p))
    return self.updates

def get_config(self):
    config = {'lr': float(K.get_value(self.lr)),
              'decay': float(K.get_value(self.decay)),
              'sigma': float(K.get_value(self.sigma))
            }
    base_config = super(DPSGD, self).get_config()
    return dict(list(base_config.items()) + list(config.items()))

```

(a)

1. Algorithm 7: Plot Training Accuracy Loss

```

dp_ax=[]
dp_ls=[]
sigs=[0.001,0.01,0.1,1.0,2.0]

for sigma in sigs:
    dp_model = build_model(dp=True,sigma=sigma)
    dp_history = dp_model.fit(x_train, one_hot_train,
                             epochs=epochs, batch_size=128)
    dp_acc = dp_history.history['acc']
    dp_loss = dp_history.history['loss']
    dp_ax.append(dp_acc)
    dp_ls.append(dp_loss)
    dp_loss_tst, dp_acc_tst =
    dp_model.evaluate(x_test, one_hot_test, verbose=1)
    print(f"DP_Model_Evaluation_after_{epochs}_epochs:\t",
          "Noise_Variance:_{sigma}\n\t\t\t\t\t",
          "Loss:_{dp_loss_tst}\n\t\t\t\t\tAccuracy:_{dp_acc_tst}")
    """

        ****Results Vary Per Iteration****
    DP Model Evaluation after 20 epochs:      Noise Variance: 0.001
                                              Loss: 0.1541813389956951
                                              Accuracy: 0.9555

    DP Model Evaluation after 20 epochs:      Noise Variance: 0.01
                                              Loss: 0.15676156164556743
                                              Accuracy: 0.9547

```

DP Model Evaluation after 20 epochs:	Noise Variance: 0.1 Loss: 0.15204105867668985 Accuracy: 0.9559
DP Model Evaluation after 20 epochs:	Noise Variance: 1.0 Loss: 3.314135622215271 Accuracy: 0.7934
DP Model Evaluation after 20 epochs:	Noise Variance: 2.0 Loss: 14.241949020385743 Accuracy: 0.1164
"""	

2.

Algorithm 8: Plot Training Accuracy Loss

3.

`plot_training(dp_ax, dp_ls, 'priv_training', sigs)`




Figure 3: Differentially Private Training Accuracy Loss, 20 Epochs**

** The plots for $\sigma=0.001, 0.01, 0.1$ overlap, as is evidenced by their accuracies at the end of training, printed above.